
Comparative Analysis of Tokenization Methods for Korean Sentiment Analysis: SentencePiece vs. Mecab

KeeTaek Yang
AIFEL online no.9

Abstract

Tokenization is a critical preprocessing step in natural language processing (NLP), especially for morphologically rich languages like Korean. In this study, I investigate the impact of two distinct tokenization methods, SentencePiece and Mecab, on sentiment analysis performance using the Naver Sentiment Movie Corpus (NSMC). SentencePiece, a subword-based tokenizer, offers flexibility by creating tokens based on statistical frequency, while Mecab, a morpheme-based tokenizer, segments text according to linguistic rules specific to Korean. I implement a bidirectional LSTM model to evaluate each tokenization method, measuring accuracy, precision, and recall. The results indicate that Mecab, with its focus on linguistic morphemes, captures sentiment-bearing nuances more effectively, outperforming SentencePiece in terms of overall accuracy. This analysis provides insights into the strengths and weaknesses of each approach, offering practical recommendations for Korean NLP tasks where sentiment nuances are crucial. Future work includes exploring hybrid tokenization approaches that balance linguistic specificity with computational efficiency.

1 Introduction to Tokenization Challenges in Korean Sentiment Analysis

The rapid advancement of natural language processing (NLP) has significantly enhanced our ability to analyze and understand text in multiple languages. However, applying these techniques to Korean, a language with unique morphological and syntactic structures, presents distinct challenges. Korean relies on complex word morphology, where a single word can express nuanced meanings through particles and affixes, and lacks clear word delimiters found in languages like English. Such characteristics pose difficulties for NLP models, making tokenization – the process of dividing text into linguistically meaningful units – a critical pre-processing step in Korean NLP tasks, including sentiment analysis.

Tokenization directly impacts model performance, as different tokenization methods can lead to variations in vocabulary size, representation accuracy, and model interpretability. For sentiment analysis in particular, where the goal is to capture emotional tone and sentiment-bearing expressions, effective tokenization can improve a model's ability to detect subtle sentiment shifts and nuanced expressions. Traditional morpheme-based tokenizers, like Mecab, have been widely used to handle Korean's rich morphology by segmenting words into their constituent morphemes. In contrast, subword methods such as SentencePiece, which are based on data-driven approaches, allow for flexible vocabulary control and address the challenges posed by out-of-vocabulary words. Despite their strengths, the relative effectiveness of these methods in capturing sentiment-bearing units in Korean text remains underexplored.

38 This study aims to address this gap by comparing the performance of SentencePiece and
 39 Mecab on a sentiment analysis task using the Naver Sentiment Movie Corpus (NSMC),
 40 a popular dataset of Korean movie reviews. By examining both quantitative metrics and
 41 qualitative case studies, I seek to understand how each tokenization method affects sentiment
 42 analysis outcomes, providing insights into their strengths and weaknesses in the Korean NLP
 43 context. Through this comparative analysis, I hope to inform NLP practitioners on the
 44 advantages and limitations of morpheme-based versus subword tokenization methods, with
 45 practical implications for Korean-language sentiment analysis and other NLP applications.

46 2 Related Work

47 2.1 Tokenization Approaches in Korean Language Processing

48 Tokenization is crucial for segmenting text into manageable units, enabling syntactic and
 49 semantic analysis in NLP tasks. Korean, characterized by agglutinative structures and
 50 complex morphology, presents unique tokenization challenges [?]. Words often comprise
 51 multiple morphemes—units carrying distinct syntactic or semantic information—making it
 52 difficult to apply tokenization methods designed for languages with simpler morphological
 53 structures.

54 Traditional Korean tokenizers like MeCab operate at the morpheme level, decomposing
 55 words into smaller, meaningful units such as nouns, particles, and verb endings [?]. This
 56 approach effectively captures syntactic nuances, allowing NLP models to accurately parse
 57 morphologically complex words. MeCab, being dictionary-based and rule-driven, is fre-
 58 quently employed in tasks like machine translation and dependency parsing, where precise
 59 segmentation of linguistic units significantly impacts performance [?].

60 Recent subword tokenization techniques, such as Byte-Pair Encoding (BPE) and Sentence-
 61 Piece, adopt a data-driven approach by generating token vocabularies based on statistical
 62 frequencies rather than predefined linguistic rules [?]. SentencePiece, in particular, creates
 63 subword units that efficiently manage vocabulary size and mitigate out-of-vocabulary issues.
 64 This flexibility makes it well-suited for deep learning applications, where vocabulary control
 65 and representation learning are crucial—especially in morphologically rich languages like
 66 Korean [?].

67 2.2 Previous Research in Korean Sentiment Analysis

68 Sentiment analysis, which involves extracting subjective information and emotional tone
 69 from text, is an active research area in Korean NLP. The Naver Sentiment Movie Corpus
 70 (NSMC) serves as a primary benchmark dataset, containing thousands of movie reviews
 71 annotated with positive or negative sentiments. It enables researchers to explore the nuances
 72 of sentiment expressions in Korean [?].

73 Studies demonstrate that tokenization quality is critical for sentiment analysis performance,
 74 as sentiment-bearing words or morphemes in Korean often convey subtle shifts in tone and
 75 intensity [?]. Korean text, especially in social media and informal reviews, includes emotive
 76 particles, abbreviations, and slang that challenge conventional tokenizers. Research high-
 77 lights how tokenization affects model accuracy by influencing the clarity and interpretability
 78 of sentiment expressions [?]. For instance, morpheme-based tokenizers like MeCab can cap-
 79 ture nuances in grammatical endings that signal emotional tone, while subword tokenizers
 80 like SentencePiece handle rare words and subwords more flexibly, addressing vocabulary
 81 limitations and improving classification performance [?].

82 Although comparative studies have examined tokenization methods for Korean language
 83 tasks, research specific to sentiment analysis remains limited. Recent findings suggest that
 84 morpheme-based tokenizers excel in tasks requiring syntactic precision, such as dependency
 85 parsing, whereas subword tokenizers generalize effectively in neural network models where
 86 vocabulary size control and representation are critical [?]. Building on these findings, this
 87 paper compares SentencePiece and MeCab in the context of sentiment analysis, providing
 88 insights into how each method affects model performance on Korean text.

89 3 Methodology and Experimental Setup

90 3.1 Dataset: Naver Sentiment Movie Corpus (NSMC)

91 For this study, I utilized the Naver Sentiment Movie Corpus (NSMC), a widely used dataset
92 for sentiment analysis in Korean. The NSMC dataset consists of thousands of Korean
93 movie reviews, each labeled with a binary sentiment (positive or negative). This corpus is
94 well-suited for tokenization experiments due to its informal language style, including slang,
95 abbreviations, and emotive expressions common in online reviews.

96 Preprocessing Steps:

- 97 • **Cleaning and Normalization:** Basic cleaning was applied to remove extraneous
98 characters and standardize text.
- 99 • **Padding and Truncation:** Sentences were padded or truncated to a fixed length
100 to ensure uniform input sizes, optimizing them for model training.

101 3.2 Tokenization Techniques: SentencePiece and Mecab

102 Two tokenization methods were applied to the NSMC dataset for comparison: SentencePiece
103 and Mecab.

104 **SentencePiece:** SentencePiece is a data-driven subword tokenizer that does not rely on
105 predefined linguistic rules. It splits text into subword units based on frequency, resulting in
106 a flexible vocabulary size and fewer out-of-vocabulary words. SentencePiece was configured
107 to generate subwords without removing punctuation, aiming to capture meaningful semantic
108 fragments within words.

109 **Mecab:** Mecab is a rule-based, morpheme-level tokenizer commonly used in Korean NLP.
110 It segments words into morphemes based on linguistic rules, capturing grammatical nuances
111 critical for syntactic interpretation. For this experiment, Mecab was configured to remove
112 punctuation to focus on semantic elements rather than sentence structure.

113 3.2.1 SentencePiece Tokenizer Setup

114 The following code snippet shows how the SentencePiece tokenizer was configured and used
115 for subword tokenization.

```
116 import sentencepiece as spm
117
118 # Train SentencePiece model
119 spm.SentencePieceTrainer.train(input='nsmc.txt', model_prefix='spm',
120                               vocab_size=32000)
121
122 # Load and tokenize using SentencePiece
123 sp = spm.SentencePieceProcessor(model_file='spm.model')
124 tokens = sp.encode_as_pieces("영화가 정말재미있었어요 !")
125 print(tokens)
```

Listing 1: SentencePiece Tokenization Example

128 **Output:** [' 영화', '가', ' 정말', ' 재미있', '었', '어요', '!']

129 3.2.2 Mecab Tokenizer Setup

130 The following code snippet demonstrates how the Mecab tokenizer was configured for
131 morpheme-based tokenization.

```
132 from konlpy.tag import Mecab
133
134 mecab = Mecab()
135 tokens = mecab.morphs("영화가 정말재미있었어요 !")
136 print(tokens)
```

139 3.3 LSTM Model Architecture

140 The following code snippet outlines the LSTM model architecture used for sentiment anal-
141 ysis.

```

142 model_lstm = tf.keras.Sequential()
143 model_lstm.add(tf.keras.layers.Embedding(vocab_size, word_vector_dim,
144     input_shape=(None,)))
145
146 # First LSTM layer (Bidirectional) - BatchNormalization, Dropout applied
147 model_lstm.add(tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64,
148     return_sequences=True)))
149 model_lstm.add(tf.keras.layers.BatchNormalization())
150 model_lstm.add(tf.keras.layers.Dropout(0.4))
151
152 # Second LSTM layer (Bidirectional) - BatchNormalization, Dropout applied
153 model_lstm.add(tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64)))
154 model_lstm.add(tf.keras.layers.BatchNormalization())
155 model_lstm.add(tf.keras.layers.Dropout(0.4))
156
157 # Dense layer with L2-normalization
158 model_lstm.add(tf.keras.layers.Dense(64, activation='relu',
159     kernel_regularizer=tf.keras.regularizers.l2(0.02)))
160 model_lstm.add(tf.keras.layers.Dropout(0.5))
161
162 # Output layer
163 model_lstm.add(tf.keras.layers.Dense(1, activation='sigmoid'))
164
165 model_lstm.summary()
166
167 
```

Listing 3: LSTM Model Architecture

168 Output:

169 Model: "sequential"

170 Layer (type)	Output Shape	Param #
171 =====	=====	=====
172 embedding (Embedding)	(None, None, 10)	100000
173 -----	-----	-----
174 bidirectional_1 (Bidirection	(None, None, 128)	38400
175 -----	-----	-----
176 batch_normalization_1 (Batch	(None, None, 128)	512
177 -----	-----	-----
178 dropout_1 (Dropout)	(None, None, 128)	0
179 -----	-----	-----
180 bidirectional_2 (Bidirection	(None, 128)	98816
181 -----	-----	-----
182 batch_normalization_2 (Batch	(None, 128)	512
183 -----	-----	-----
184 dropout_2 (Dropout)	(None, 128)	0
185 -----	-----	-----
186 dense_1 (Dense)	(None, 64)	8256
187 -----	-----	-----
188 dropout_3 (Dropout)	(None, 64)	0
189 -----	-----	-----
190 dense_2 (Dense)	(None, 1)	65
191 =====	=====	=====
192		

193 Total params: 246,561
 194 Trainable params: 246,049
 195 Non-trainable params: 512
 196 -----

197 3.4 Evaluation Setup and Metrics

198 The sentiment analysis model used for both tokenization approaches was a Bidirectional
 199 LSTM (BiLSTM) model, optimized to capture context from both directions in the text.

- 200 • **Model Architecture:** The architecture included an embedding layer, two BiLSTM
 201 layers, batch normalization, dropout layers for regularization, and dense layers for
 202 classification.
- 203 • **Hyperparameters:** The BiLSTM model was trained for 20 epochs with a batch
 204 size of 1024. Dropout rates were set to 0.4 for the LSTM layer and 0.5 for the dense
 205 layer to mitigate overfitting.
- 206 • **Evaluation Metrics:** Performance was measured using accuracy, precision, and
 207 recall, providing a comprehensive evaluation of each tokenizer’s impact.

208 4 Experiments and Results: Tokenization Impact on Sentiment 209 Analysis

210 4.1 Tokenization Analysis

211 I analyzed each tokenizer’s vocabulary size, tokenization speed, and average token length:

- 212 • **Vocabulary Size:** SentencePiece created a vocabulary of subwords, which con-
 213 trolled the vocabulary size effectively. Mecab, on the other hand, generated a more
 214 extensive vocabulary based on morphemes, capturing syntactic variations.
- 215 • **Tokenization Speed:** SentencePiece demonstrated faster tokenization due to its
 216 subword-based approach, whereas Mecab, being morpheme-based, took slightly
 217 longer.
- 218 • **Token Length:** SentencePiece tokens were shorter on average, reflecting its sub-
 219 word structure, while Mecab tokens were longer due to morpheme segmentation,
 220 which resulted in fewer tokens per sentence.

221 4.2 Sentiment Model Performance

222 The table below summarizes the performance of the sentiment analysis model for each
 223 tokenization method in terms of accuracy, precision, and recall.

Table 1: Sentiment Model Performance with SentencePiece and Mecab

Tokenization Method	Criteria for Vocabulary	Accuracy	Precision	Recall
SentencePiece	Controlled subword vocabulary	0.8285	0.8078	0.8654
Mecab	Extensive morpheme-based vocabulary	0.8422	0.8484	0.8355

224 4.3 Qualitative Observations

225 Examples from the tokenized text revealed that SentencePiece captured frequent subwords,
 226 which helped in generalizing across vocabulary. However, it sometimes missed sentiment-
 227 bearing morphemes. Mecab, with its morpheme-level granularity, was better at capturing
 228 sentiment nuances specific to the Korean language, such as mood particles and verb endings,
 229 which contribute to the emotional tone of sentences.

230 5 Discussion: Insights on Tokenization for Korean NLP

231 5.1 Interpreting Results in Korean Language Context

232 The experiments indicate that tokenization method influences sentiment analysis perfor-
233 mance. Mecab, with its focus on morphemes, outperformed SentencePiece in this study,
234 likely due to its ability to capture Korean-specific syntactic features that influence senti-
235 ment. SentencePiece, while efficient, may lose sentiment nuances by segmenting words into
236 subwords without considering linguistic rules.

237 5.2 Strengths and weaknesses of SentencePiece vs. Mecab

238 **SentencePiece:** Its data-driven approach makes it versatile and computationally efficient,
239 suitable for applications requiring fast processing and a manageable vocabulary. however, it
240 may lack the linguistic specificity needed for tasks that depend on syntactic or morphological
241 cues, such as sentiment analysis in Korean.

242 **Mecab:** While slower and more complex, Mecab’s rule-based segmentation aligns well with
243 the structural nuances of Korean, making it more effective for sentiment-sensitive tasks.
244 Its reliance on linguistic rules allows it to capture important morphemes, although it may
245 require more computational resources.

246 5.3 Broader Implications for Korean NLP

247 These findings suggest that morpheme-based tokenization methods like Mecab may be prefer-
248 able for sentiment analysis and other Korean NLP tasks that depend on syntactic or emotive
249 nuances. however, subword tokenizers like SentencePiece could be more efficient for appli-
250 cations requiring generalization across diverse domains and less dependency on language
251 structure.

252 6 Conclusion and Future Directions

253 This study compared two tokenization methods—SentencePiece and Mecab—on a sentiment
254 analysis task using the Naver Sentiment Movie Corpus. The results shoId that Mecab’s
255 morpheme-based approach was better suited for capturing sentiment in Korean, while Sen-
256 tencePiece offered faster processing and flexibility at the cost of some sentiment-specific
257 detail.

258 **Future Work:** Further studies could investigate hybrid tokenization methods that combine
259 subword and morpheme-based approaches to balance speed and linguistic precision. Addi-
260 tionally, experiments on other Korean NLP tasks, such as translation or summarization,
261 could offer broader insights into tokenization strategies for Korean text processing.

262 References

263 A Appendix